

Recommender systems

Matrix Factorization models

Lecture 3
2025

Anna Volodkevich

Contents

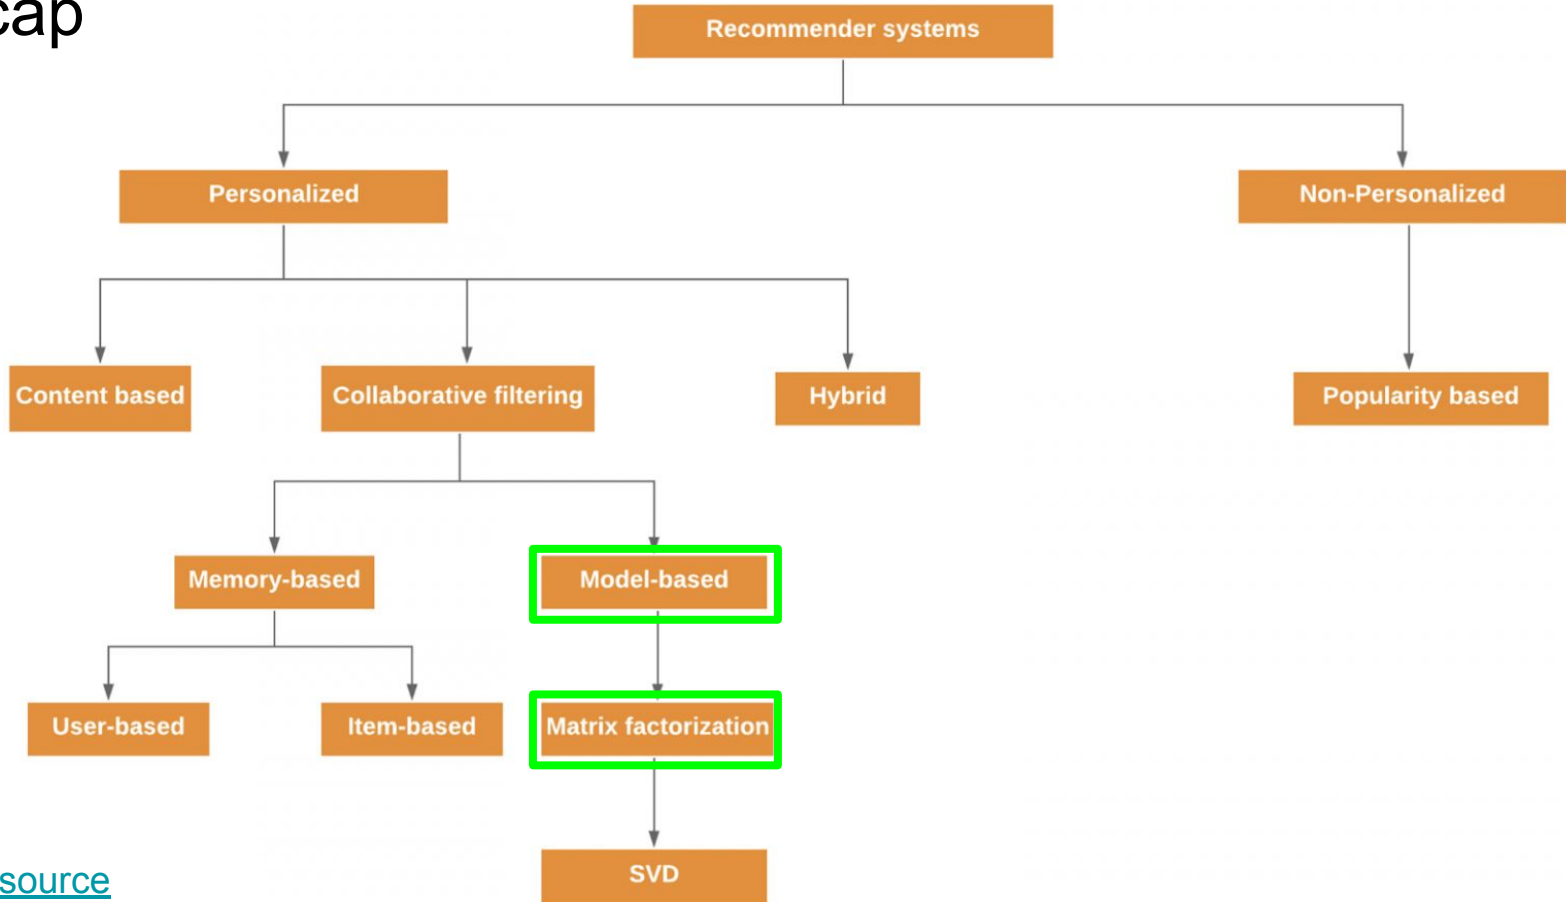
RS as a matrix completion task

SVD, Truncated SVD

Funk SVD

ALS, Implicit ALS

Recap



Recap

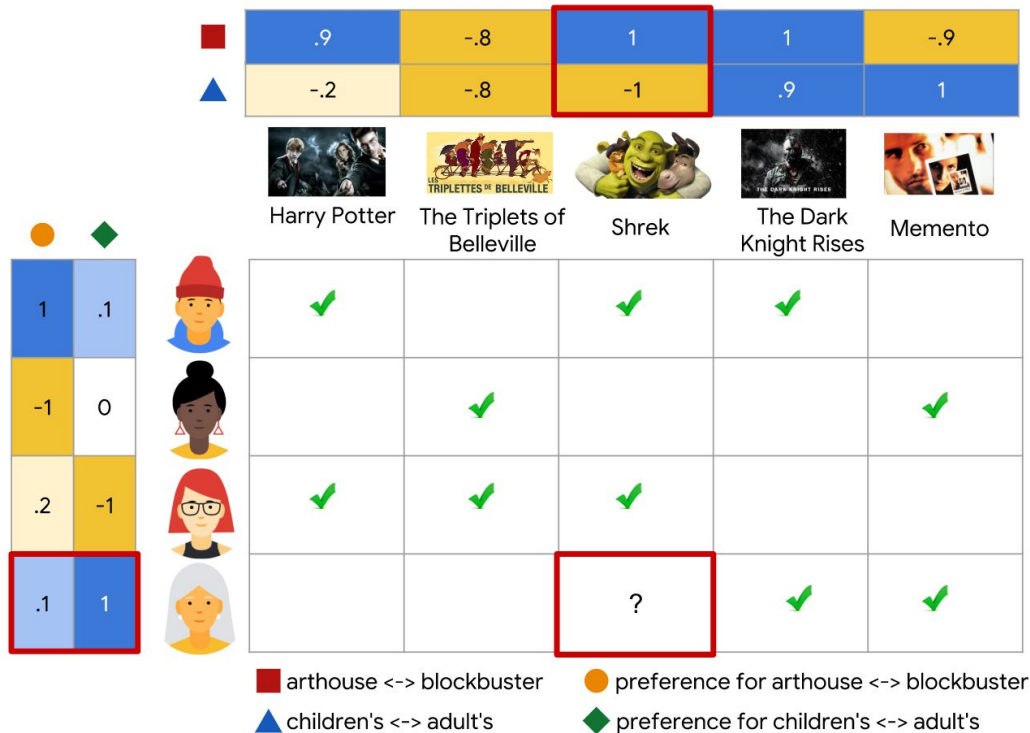
Model-based methods

In model-based methods, machine learning and data mining methods are used in the context of predictive models. In cases where the model is parameterized, the parameters of this model are learned within the context of an optimization framework. Some examples of such model-based methods include decision trees, rule-based models, Bayesian methods and **latent factor models**. Many of these methods, such as latent factor models, have a high level of coverage even for sparse ratings matrices.

Latent factor models, such as matrix factorization, transform both items and users to the same latent factor space. The latent space tries to explain ratings by characterizing both products and users on factors automatically inferred from user feedback.

Recommendations as a matrix completion

- User-item interactions are represented in a matrix form
- We assume that the users and items could be represented in a low-dimensional latent feature space
- MNAR problem
- Need to consider different approaches for different feedback types



[Image source](#)

SVD: recap

- Σ - diagonal $m \times n$ matrix of non-increasing singular values $\sigma_1, \dots, \sigma_n$
- U и V – orthogonal matrices left-singular vectors and right-singular vectors

The diagram illustrates the SVD decomposition of a matrix M into three matrices: U , Σ , and V^* .

Top Row:

- M (4x4): A 4x4 grid of gray squares.
- U (4x4): A 4x4 grid with columns colored teal, green, blue, and green.
- Σ (4x4): A 4x4 grid with diagonal elements colored orange, yellow, and yellow, and all other elements gray.
- V^* (4x4): A 4x4 grid with rows colored light blue, purple, purple, and pink.

Equation: $M = U \Sigma V^*$

Dimensions: $m \times n$ $m \times m$ $m \times n$ $n \times n$

Bottom Row:

- U (4x4): A 4x4 grid with columns colored teal, green, blue, and green.
- U^* (4x4): A 4x4 grid with rows colored teal, green, blue, and green.
- I_m (4x4): A 4x4 grid with diagonal elements colored teal, green, blue, and green, and all other elements gray.


Equation: $U U^* = I_m$

Bottom Row:

- V (4x4): A 4x4 grid with columns colored light blue, purple, and pink.
- V^* (4x4): A 4x4 grid with rows colored light blue, purple, and pink.
- I_n (4x4): A 4x4 grid with diagonal elements colored light blue, purple, and pink, and all other elements gray.

Equation: $V V^* = I_n$

Truncated SVD



$$R \approx U_k \Sigma_k V_k^T$$

$$R \approx M V_k V_k^T$$

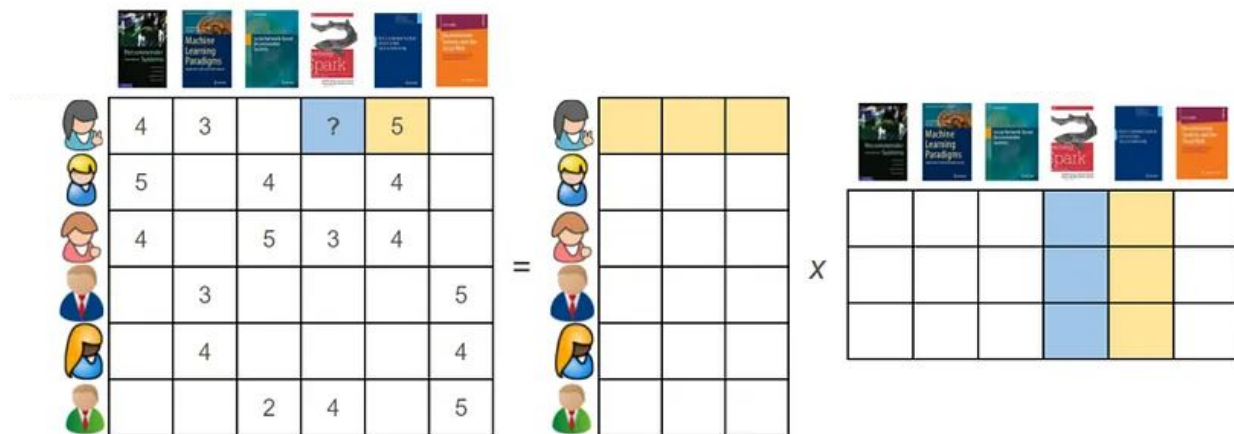
[SVD visualization](#)

$$\|R - U_k \Sigma_k V_k^t\|_F^2 \rightarrow \min$$

- Could be interpreted as calculating user-item similarity in latent features space considering top-k most important topics
- The best top-k approximation in terms of Frobenius norm
- Folding-in and User/item vector update without retraining

Funk SVD

$$R \approx PQ^T$$




$$\sum_{(i,j) \in R} (r_{ij} - p_i^T q_j)^2 + \sum_{i=1}^m \lambda_1 ||p_i||^2 + \sum_{j=1}^n \lambda_2 ||q_j||^2 \rightarrow \min$$

[Image source](#)

Similarity in latent space

- MF provides a solution for top-k personalized recommendations and item2item recommendations
- approximate nearest neighbours search could be applied for fast recommendation building



[Image source](#)

Recup: cosine similarity, dot product

Dot product

$$\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\| \|\mathbf{b}\| \cos \theta,$$

where θ is the **angle** between \mathbf{a} and \mathbf{b} .


\mathbf{a} is projected to \mathbf{b} :

$$a_b = \|\mathbf{a}\| \cos \theta,$$

$$\mathbf{a} \cdot \mathbf{b} = a_b \|\mathbf{b}\| = b_a \|\mathbf{a}\|.$$

Cosine similarity

$$S_C(A, B) := \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|}$$



SVD and Funk SVD: pros and cons

Disadvantages

- do not use features and context (could be extended, see Frolov, E., & Oseledets, I. (2019, September). HybridSVD: when collaborative information is not enough)
- Truncated SVD: consider empty values as zeros
- Funk SVD: do not consider empty values
- popularity bias

Advantages

- strong baseline out of the box ([scipy](#), [sklearn](#))
- fast inference with [ANN](#)
- Truncated SVD: new user/item vectors inference with closed-form solution (folding in)
- Funk SVD learning objective could be customized for a particular task

Funk SVD customization

- Adding biases

$$\hat{r}_{ui} = \mu + b_i + b_u + q_i^T p_u.$$

$$\min_{b_*, q_*, p_*} \sum_{(u,i) \in \mathcal{K}} (r_{ui} - \mu - b_i - b_u - q_i^T p_u)^2 + \lambda_4 (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2).$$

- SVD++: adding implicit feedback

$$\hat{r}_{ui} = \mu + b_i + b_u + q_i^T \left(p_u + |\mathcal{R}(u)|^{-\frac{1}{2}} \sum_{j \in \mathcal{R}(u)} y_j \right)$$

y_j - implicit feedback item vectors,
 q_i - explicit feedback item vector

- timeSVD, timeSVD++: adding time dependencies
- ...

Alternating least squares

$$\min_{X,Y} \sum_{r_{ui} \text{ observed}} (r_{ui} - x_u^\top y_i)^2 + \lambda \left(\sum_u \|x_u\|^2 + \sum_i \|y_i\|^2 \right)$$

Our approach will be to fix Y and optimize X , then fix X and optimize Y , and repeat until convergence.

Alternating least squares

Algorithm 1 ALS for Matrix Completion

Initialize X, Y

repeat

for $u = 1 \dots n$ **do**

$$x_u = \left(\sum_{r_{ui} \in r_{u*}} y_i y_i^\top + \lambda I_k \right)^{-1} \sum_{r_{ui} \in r_{u*}} r_{ui} y_i \quad (2)$$

end for

for $i = 1 \dots m$ **do**

$$y_i = \left(\sum_{r_{ui} \in r_{*i}} x_u x_u^\top + \lambda I_k \right)^{-1} \sum_{r_{ui} \in r_{*i}} r_{ui} x_u \quad (3)$$

end for

until convergence

Alternating least squares for implicit feedback (iALS)

$$p_{ui} = \begin{cases} 1 & r_{ui} > 0 \\ 0 & r_{ui} = 0 \end{cases} \quad c_{ui} = 1 + \alpha r_{ui}$$

p_{ui} - feedback presence (bool)
 c_{ui} - level of certainty

$$\min_{x_*, y_*} \sum_{u, i} c_{ui} (p_{ui} - x_u^T y_i)^2 + \lambda \left(\sum_u \|x_u\|^2 + \sum_i \|y_i\|^2 \right)$$

$f \times f$ matrix $Y^T Y$ in time $O(f^2 n)$. For each user u , let us define the diagonal $n \times n$ matrix C^u where $C_{ii}^u = c_{ui}$, and also the vector $p(u) \in \mathbb{R}^n$ that contains all the preferences by u (the p_{ui} values). By differentiation we find an analytic expression for x_u that minimizes the cost function (3):

$$x_u = (Y^T C^u Y + \lambda I)^{-1} Y^T C^u p(u) \quad \rightarrow \quad x_u = (Y^T Y + \lambda I + \sum_{\forall i | p_{ui} \neq 0} (c_{ui} - 1) y_i y_i^T)^{-1} \left(\sum_{\forall i | p_{ui} \neq 0} c_{ui} p_{ui} y_i \right)$$

$f \times 1$ $f \times f$ $f \times n$ $f \times f$ $f \times f$

ALS/iALS: pros and cons

Disadvantages

- interactions data only (do not use features and context)
- **quadratic (MSE) loss only**

Advantages

- strong baseline out of the box ([implicit](#), [pyspark mllib](#))
- fast inference with [ANN](#)
- new user/item vectors inference with closed-form solution (folding in)
- **any feedback type**
- **good parallelisation**

Additional resources

Articles (RU)

[Habr: Рекомендательные системы: идеи, подходы, задачи](#)

[Habr: Как работают рекомендательные системы. Лекция в Яндексе](#)

[Лекция по рекомендательным системам из курса ВШЭ по машинному обучению](#)

Books

“Recommender Systems. The Textbook”, 2016, Springer

“Recommender Systems Handbook”, 2011, Springer

Frameworks

[implicit](#): iALS and item-based KNN