

Recommender systems

# Multi-stage recommender systems

Lecture 5  
Fall 2023

Anna Volodkevich

# Contents

General multi-stage recommendation pipeline

Data split

First-stage models

Feature generation

Second-stage models

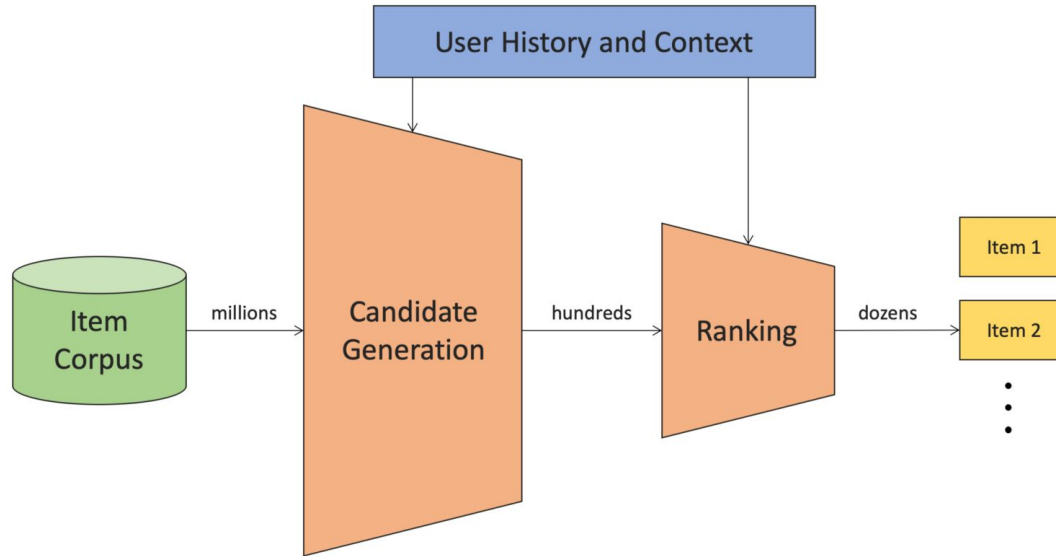
# Recap: recommendation as a general machine learning task on tabular data

Feature vector $\mathbf{x}$																	Target $y$					
$\mathbf{x}^{(1)}$	1	0	0	...	1	0	0	0	...	0.3	0.3	0.3	0	...	13	0	0	0	0	...	5	$y^{(1)}$
$\mathbf{x}^{(2)}$	1	0	0	...	0	1	0	0	...	0.3	0.3	0.3	0	...	14	1	0	0	0	...	3	$y^{(2)}$
$\mathbf{x}^{(3)}$	1	0	0	...	0	0	1	0	...	0.3	0.3	0.3	0	...	16	0	1	0	0	...	1	$y^{(2)}$
$\mathbf{x}^{(4)}$	0	1	0	...	0	0	1	0	...	0	0	0.5	0.5	...	5	0	0	0	0	...	4	$y^{(3)}$
$\mathbf{x}^{(5)}$	0	1	0	...	0	0	0	1	...	0	0	0.5	0.5	...	8	0	0	1	0	...	5	$y^{(4)}$
$\mathbf{x}^{(6)}$	0	0	1	...	1	0	0	0	...	0.5	0	0.5	0	...	9	0	0	0	0	...	1	$y^{(5)}$
$\mathbf{x}^{(7)}$	0	0	1	...	0	0	1	0	...	0.5	0	0.5	0	...	12	1	0	0	0	...	5	$y^{(6)}$
A B C ... User				TI NH SW ST ... Movie					TI NH SW ST ... Other Movies rated					Time	TI NH SW ST ... Last Movie rated							

## Conditions:

- We want to use all available features of different types and nature
- Some features depend on user-item pair and thus should be calculated online
- Negatives (not relevant) are often not available
- It is impractical to score all user-item pairs with tabular models
- Production models should be fast
- Possible solution: hand-crafted negative generation and **multi-stage recommender systems**

# Multi-stage recommenders: Reasoning and Pipeline



- collaborative, content-based, non-personalized and the other models has their own advantages and disadvantages
- developers want to **combine recommendations** from different sources
- let's generate a small set of candidates for each user with light and fast models and **rerank** them with the complex one **to get top-k**

# General multi-stage recommendation pipeline from Nvidia

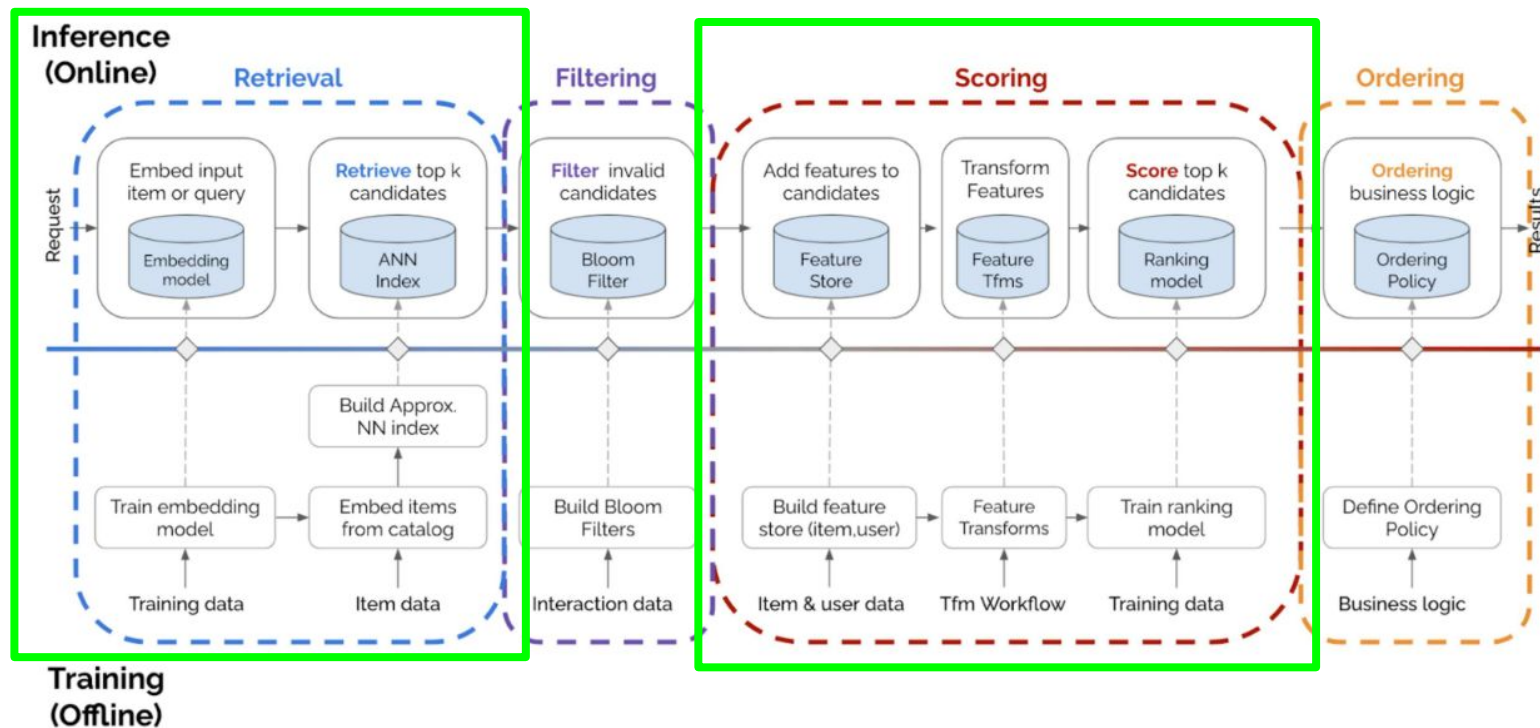
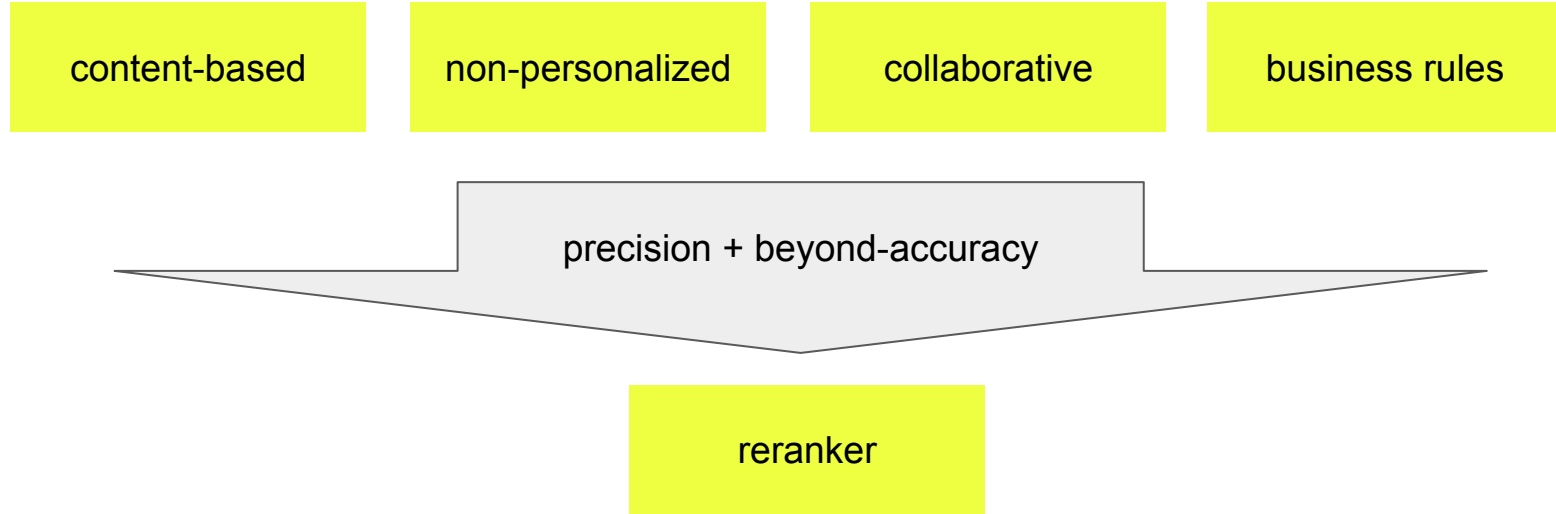


Figure 1: An overview of Four-stage Recommender Systems

[Image source](#)

# Two-stage model candidate sources



Tunable parameters / choices:

- number of candidates from each model
- first-stage model types

# Two-stage model feature sources

Item/user features  
Context features

First-level outputs

- scores
- ranks
- vectors and biases

User/ item  
statistics


- popularity / ctr over time
- popularity tendency
- time features
- target encoding

Co-occurrences and pair features

- number of interactions with an item/category in the past
- mean/popularity conditioned on feature value (popularity of an item in an age group)
- time features conditioned on feature value (last time bought an item from category)
- co-occurrence with the other items from user history

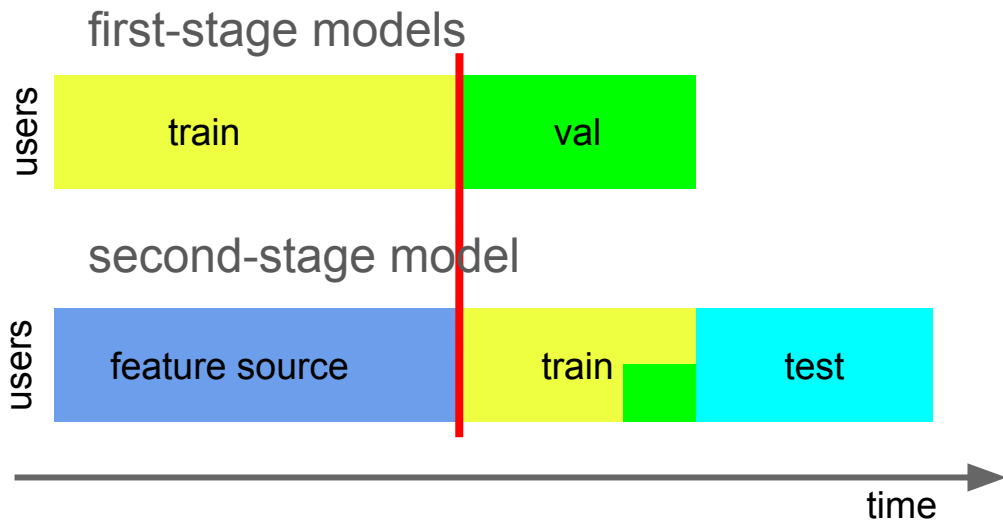
# Data split

## Simple recommender



- Want to capture the most recent patterns
- Need to avoid overfitting of the second-stage model

## Multi-stage recommender



Tunable parameters / choices:

- split ratio
- retraining (attention to possibly unstable vectors and scores)



## Second-stage pipeline

- Combine candidates from all first-stage models
- Get scores and ranks from all first-stage models
- Add negative examples and target (usually 0 and 1)
- Add features
- Train the model **or**
- Make prediction, get top-k and calculate metric

Tunable parameters / choices:

- source of negative items: first stage, random,  $\propto$  popularity
- second-stage model type, task and loss function

# Deep learning as a promising alternative

**For the past few years most published research on recommendation algorithms has been based on deep learning (DL) methods.** Following common research practices in our field, these works usually demonstrate that a **new DL method is outperforming other models not based on deep learning in offline experiments.** This almost consistent success of DL based models is however not observed in recommendation-related machine learning competitions like the challenges that are held with the yearly ACM RecSys conference. **Instead the winning solutions mostly consist of substantial feature engineering efforts and the use of gradient boosting or ensemble techniques.**

*Why Are Deep Learning Models Not Consistently Winning Recommender Systems Competitions Yet?: A Position Paper RecSys'20*

## Deep learning models

- can utilize all previously discussed gradient-based models
- can learn embeddings and work with all data types end-to-end (no separate models retraining)
- provide simple sequential data modelling

# Additional sources

## Multi-stage models

[Your second recsys: Multi-stage models](#)

[Rekko Challenge: second place solution](#)

[RecSys Challenge 2018: third place solution](#)

[Dzen Recommendation system](#)

[OKKO Recommendation system](#)

## Other interesting papers and videos

Meetup series “Дзен-митап” (youtube)

[Yandex Music Recommendation System](#)