

Recommender systems

Neural recommenders

Lecture 6
2024

Anna Volodkevich

Contents

MF as a neural network

Negative sampling

Incorporating features: Two-tower/DSSM

Recap: ranking loss functions

Adding feature cross (DCN-V2)

Sequential models: SASRec

Matrix factorization as a neural network

MF formulation: Recap

MF associates each user and item with a real-valued vector of latent features. Let \mathbf{p}_u and \mathbf{q}_i denote the latent vector for user u and item i , respectively; MF estimates an interaction y_{ui} as the inner product of \mathbf{p}_u and \mathbf{q}_i :

$$\hat{y}_{ui} = f(u, i | \mathbf{p}_u, \mathbf{q}_i) = \mathbf{p}_u^T \mathbf{q}_i = \sum_{k=1}^K p_{uk} q_{ik}, \quad (2)$$

Regression formulation

To learn model parameters, existing pointwise methods [14, 39] largely perform a regression with squared loss:

$$L_{sq} = \sum_{(u,i) \in \mathcal{Y} \cup \mathcal{Y}^-} w_{ui} (y_{ui} - \hat{y}_{ui})^2, \quad (5)$$

where \mathcal{Y} denotes the set of observed interactions in \mathbf{Y} , and \mathcal{Y}^- denotes the set of negative instances, which can be all (or sampled from) unobserved interactions; and w_{ui} is a hyper-parameter denoting the weight of training instance (u, i) .

Classification formulation

$$p(\mathcal{Y}, \mathcal{Y}^- | \mathbf{P}, \mathbf{Q}, \Theta_f) = \prod_{(u,i) \in \mathcal{Y}} \hat{y}_{ui} \prod_{(u,j) \in \mathcal{Y}^-} (1 - \hat{y}_{uj}). \quad (6)$$

Taking the negative logarithm of the likelihood, we reach

$$\begin{aligned} L &= - \sum_{(u,i) \in \mathcal{Y}} \log \hat{y}_{ui} - \sum_{(u,j) \in \mathcal{Y}^-} \log(1 - \hat{y}_{uj}) \\ &= - \sum_{(u,i) \in \mathcal{Y} \cup \mathcal{Y}^-} y_{ui} \log \hat{y}_{ui} + (1 - y_{ui}) \log(1 - \hat{y}_{ui}). \end{aligned} \quad (7)$$

Negative sampling

- **Uniform negative sampling.** For each training example (consisting of a user and a positive item), we sample m negative (unobserved) items, uniformly at random. May be slow to converge (“easy” negatives).
- **Popularity-based negative sampling.** For each training example (consisting of a user and a positive item), we sample m negative (unobserved) items, proportionally to their popularity in dataset.
- **“Hard” negative sampling.** Sample items using previous model, giving the maximal gradient update by being similar to the user’s positives or giving the highest score.
- **Rule-based negatives.** E.g. negatives from the same category.
- **In-batch negatives.** The other users positives are the user negatives.

More on negative sampling

- [Chen, Chong, et al. "Revisiting negative sampling vs. non-sampling in implicit recommendation." ACM Transactions on Information Systems 41.1 \(2023\): 1-25.](#)
- [Pellegrini, Roberto, Wenjie Zhao, and Iain Murray. "Don't recommend the obvious: estimate probability ratios." Proceedings of the 16th ACM Conference on Recommender Systems. 2022.](#)
- [Wang, HaiYing, Anan Zhang, and Chong Wang. "Nonuniform negative sampling and log odds correction with rare events data." Advances in Neural Information Processing Systems 34 \(2021\): 19847-19859.](#)

Matrix factorization as a neural network





Figure 1: A model with dot product similarity (left) and MLP-based learned similarity (right).

Matrix factorization as a neural network: NeuMF (MLP+GMF)

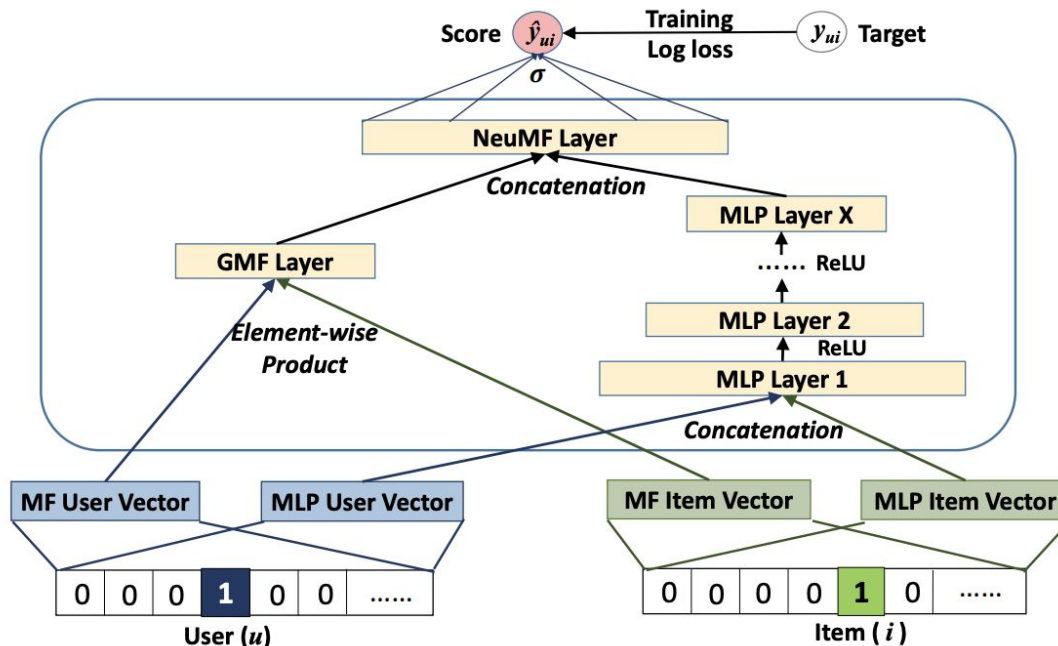


Figure 3: Neural matrix factorization model

Matrix factorization as a neural network




Figure 2: Comparison of learned similarities (MLP, NeuMF) to a dot product: The results for MLP and NeuMF are from [17]. The dot product substantially outperforms the learned similarity measures. Only the pretrained NeuMF is competitive, on one dataset, and for large embedding dimension.

Method	Movielens		Pinterest		Result from
	HR@10	NDCG@10	HR@10	NDCG@10	
Popularity	0.4535	0.2543	0.2740	0.1409	[8]
SLIM [25, 30]	<u>0.7162</u>	<u>0.4468</u>	0.8679	<u>0.5601</u>	[8]
iALS [20]	0.7111	0.4383	0.8762	0.5590	[8]
NeuMF (MLP+GMF) [17]	0.7093	0.4349	<u>0.8777</u>	0.5576	[8]
Matrix Factorization	0.7294	0.4523	0.8895	0.5794	Fig. 2

- a simple dot product substantially outperforms the proposed learned similarities
- a MLP can in theory approximate any function, we show that it is non-trivial to learn a dot product with an MLP
- feature interactions should be directly modelled

Matrix factorization as a neural network

- Advantages

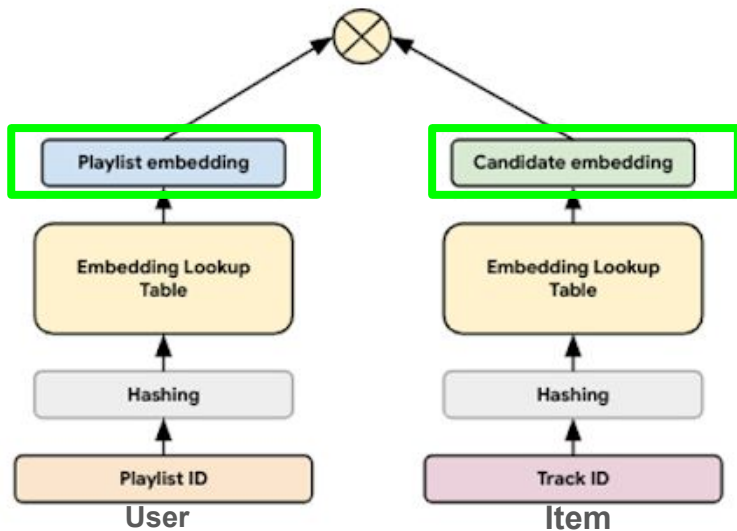
- commonly-used frameworks and optimization algorithms
- easy incremental updates
- gives user and item vectors for fast inference

- Disadvantages

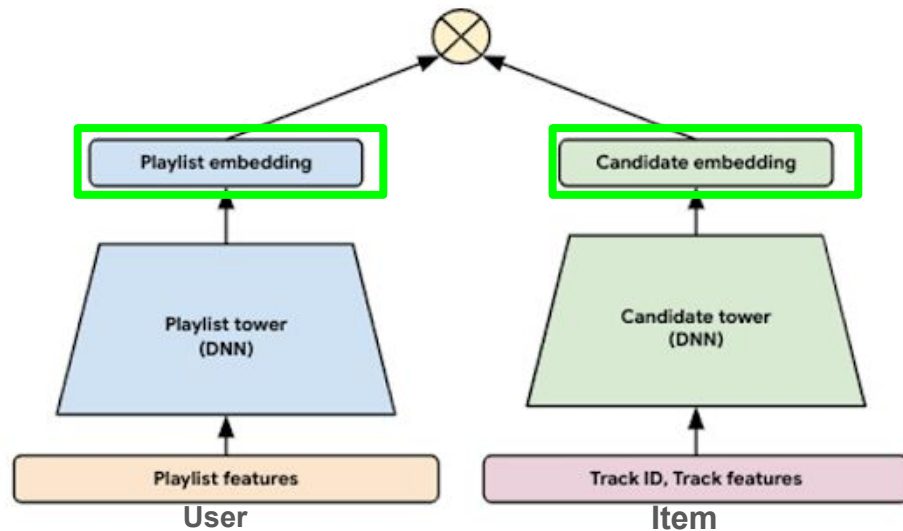
- no user/item features
- no context features
- no user sequence / latest interests representation

Incorporating features: Two-tower/DSSM

MF



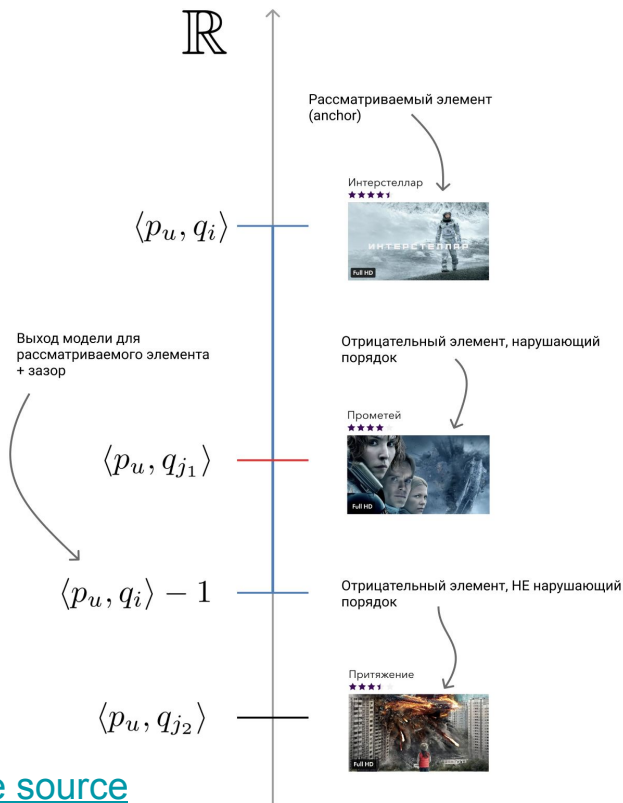
DSSM



[Huang, Po-Sen, et al. "Learning deep structured semantic models for web search using clickthrough data." Proceedings of the 22nd ACM international conference on Information & Knowledge Management. 2013.](#)

[image source: Scaling deep retrieval with TensorFlow Recommenders and Vertex AI Matching Engine](#)

Recap: ranking loss functions



[image source](#)

BPR loss

score for positive item

$$\hat{x}_{uij} := \hat{x}_{ui} - \hat{x}_{uj}$$

score for negative item

$$\sigma(x) := \frac{1}{1 + e^{-x}}$$

$$\text{BPR-OPT} := \sum_{(u,i,j) \in D_S} \ln \sigma(\hat{x}_{uij}) - \lambda_{\Theta} \|\Theta\|^2 \rightarrow \max$$

model params

Margin loss / WARP loss

score for positive item

score for negative item

$$L(\lfloor \frac{Y-1}{N} \rfloor) |1 - f_y(x_i) + f_{\bar{y}}(x_i)|_+ \rightarrow \min$$

$$L(k) = \log(k)$$

N - number of negatives sampled
 Y - number of items

[BPR paper](#)

[WARP paper](#)

Two-tower/DSSM: feature generation

- Categorical features
 - embeddings
- Text
 - pre-trained embeddings combination
 - sentence embeddings
- Numerical
 - squashing (sigmoid over scaled feature values with learned scale and bias)
 - discretization (binning)
 - piecewise linear encoding
- Sequences
 - Long/short term user history as an additional input (user as a combination of item embeddings)




Figure 1: The piecewise linear encoding (PLE) in action for $T = 4$ (see Equation 1).

Two-tower / DSSM

- Advantages

- commonly-used frameworks and optimization algorithms
- easy incremental updates
- gives user and item vectors for fast inference
- utilize user/item features

- Disadvantages

- no user sequence / latest interests representation in original architecture
- no feature cross (like in Factorization Machines)

- Useful links

- MTS Your second recsys: [video](#), [code](#)

Adding feature cross (DCN-V2)

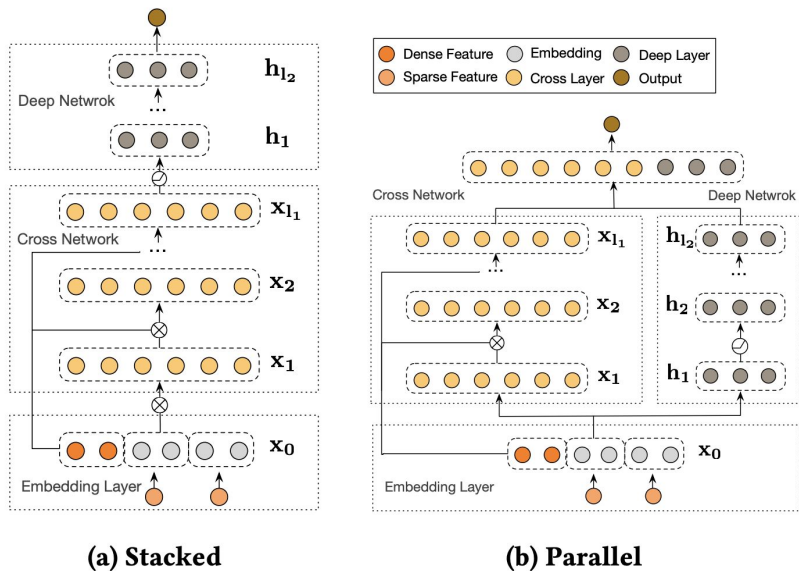


Figure 1: Visualization of DCN-V2. \otimes represents the cross operation in Eq. (1), i.e., $x_{l+1} = x_0 \odot (W_l x_l + b_l) + x_l$.

$$x_{i+1} = x_0 \odot (W \times x_i + b) + x_i$$

Figure 2: Visualization of a cross layer.

$$x_{l+1} = x_0 \odot (W_l x_l + b_l) + x_l \quad (1)$$

where $x_0 \in \mathbb{R}^d$ is the base layer that contains the original features of order 1, and is normally set as the embedding (input) layer. $x_l, x_{l+1} \in \mathbb{R}^d$, respectively, represents the input and output of the $(l+1)$ -th cross layer. $W_l \in \mathbb{R}^{d \times d}$ and $b_l \in \mathbb{R}^d$ are the learned weight matrix and bias vector. Figure 2 shows how an individual cross layer functions.

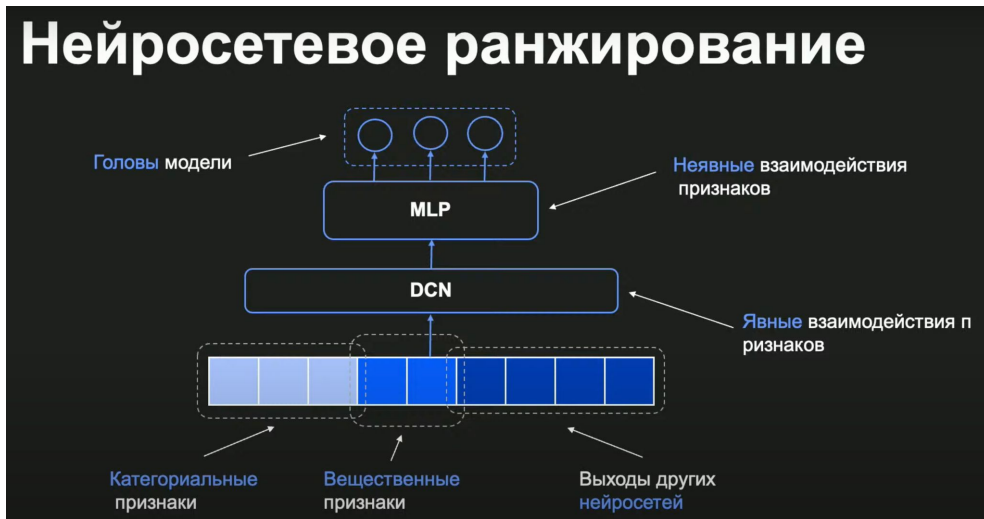
[Wang, Ruoxi. et al. "Deep & cross network for ad click predictions." Proceedings of the ADKDD'17. 2017. 1-7.](#)

[Wang, Ruoxi. et al. "Dcn v2: Improved deep & cross network and practical lessons for web-scale learning to rank systems." Proceedings of the web conference 2021. 2021.](#)

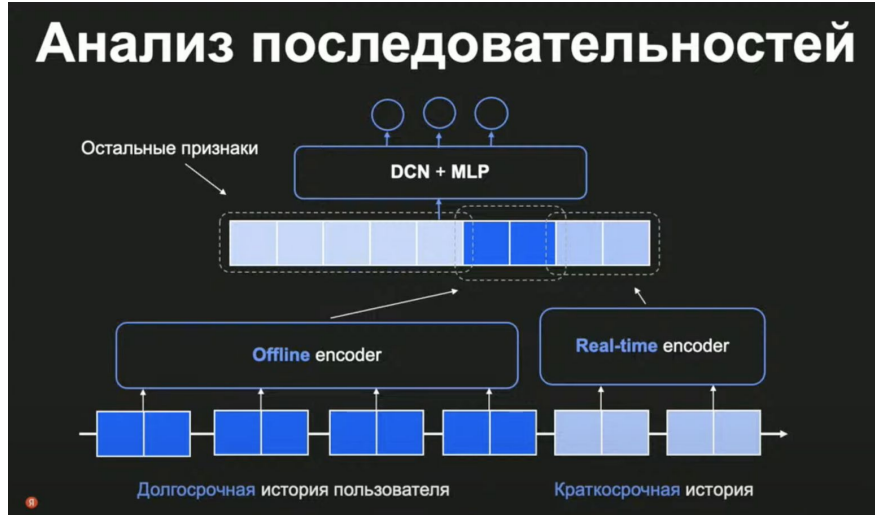
[TensorFlow Recommenders tutorial](#)

Neural ranking by Yandex

Нейросетевое ранжирование




Анализ последовательностей




- Neural ranking works for industrial tasks with large data volumes
- Allows to use more data and feature sources and gives bigger `model capacity`
- Suits for multiple feedback types out of the box

Sequential recommendations: task

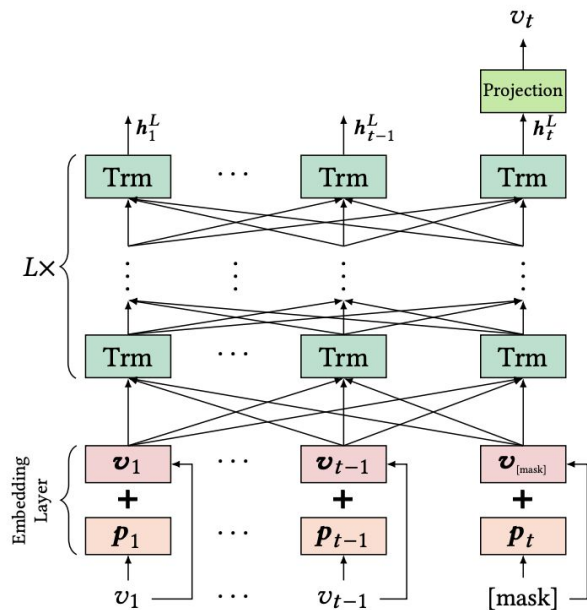


- items are treated as tokens
- users are treated as item sequences
- item embeddings are learned from user sequences with masked language modeling / next item prediction / other task
- can capture temporal dynamics and user preference drift
- suits for online recommendations
- rapidly developing research area
- commonly-used models
 - [Caser](#) (Convolutional)
 - [GRU4Rec](#)
 - [BERT4Rec](#)
 - [SASRec](#)


Sequential recommendations: examples




(a) Transformer Layer.



(b) BERT4Rec model architecture.



(c) SASRec model architecture.



(d) RNN based sequential recommendation methods.

Sequential modelling: SASRec

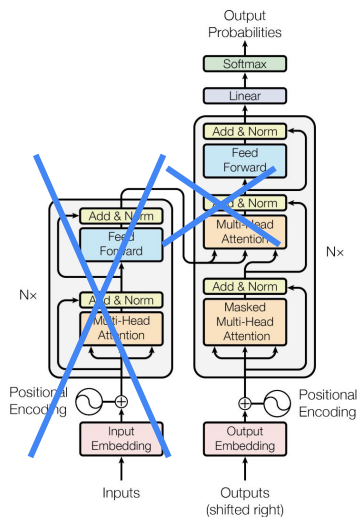
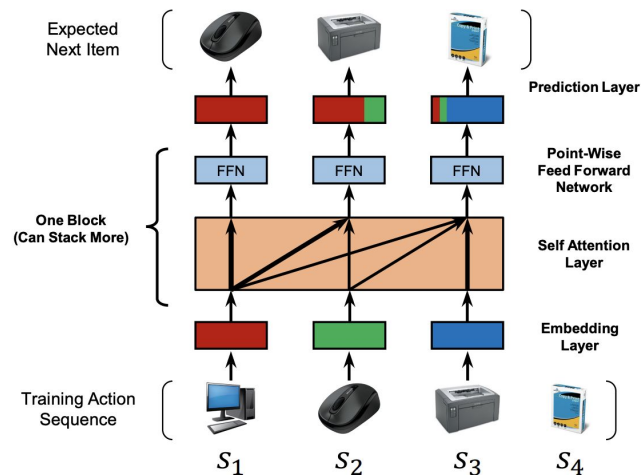




Figure 1: The Transformer - model architecture.

SASRec is a decoder-only model pretty close to language models like GPT-2

Scaled Dot-Product Attention



Multi-Head Attention



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$


- Learned positional embeddings
- [Layer normalization](#)
- Small in comparison with LMs (two blocks, two heads, latent dimensions up to 50)

[Kang, Wang-Cheng, and Julian McAuley. "Self-attentive sequential recommendation." IEEE, 2018.](#)

[Vaswani, Ashish, et al. "Attention is all you need." Advances in neural information processing systems 30 \(2017\).](#)

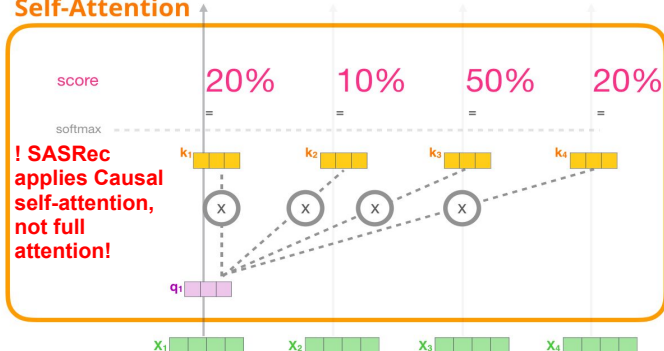
Self-Attention: Recap

Self-Attention



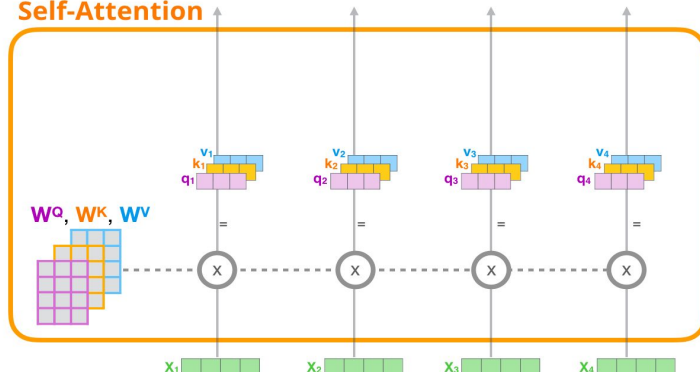
2) Multiply (dot product) the current **query vector**, by all the **key vectors**, to get a score of how well they match

Self-Attention



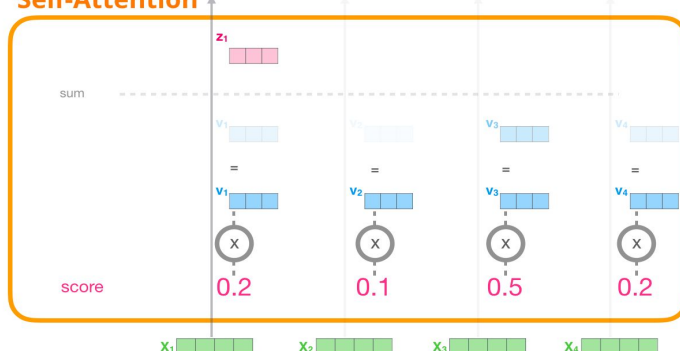
1) For each input token, create a **query vector**, a **key vector**, and a **value vector** by multiplying by weight Matrices W^q , W^k , W^v

Self-Attention



3) Multiply the **value vectors** by the **scores**, then sum up

Self-Attention



Sequential modelling: Loss functions

Original SASRec loss: binary cross entropy with one negative sample for each positive

$$\mathcal{L}_{BCE} = - \sum_{u \in U} \sum_{t=1}^{n_u} \log(\sigma(r_{t,i_t}^{(u)})) + \log(1 - \sigma(r_{t,-}^{(u)})),$$

BERT4Rec loss: full cross entropy

$$\mathcal{L}_{CE} = - \sum_{u \in U} \sum_{t \in T_u} \log \frac{\exp(r_{t,i_t}^{(u)})}{\sum_{i \in I} \exp(r_{t,i}^{(u)})}$$

Sampled cross-entropy from “Turning Dross Into Gold Loss: is BERT4Rec really better than SASRec?”

$$\mathcal{L}_{CE-sampled_N} = - \sum_{u \in U} \sum_{t=1}^{n_u} \log \frac{\exp(r_{t,i_t}^{(u)})}{\exp(r_{t,i_t}^{(u)}) + \sum_{i \in I_N^-(u)} \exp(r_{t,i}^{(u)})},$$

Sequential modelling: Open questions and research directions

- Application for different recommendation tasks, e.g. for top-k prediction task
- Quality improvement: positional encoding, negative sampling, negative feedback/different feedback types incorporation, hyperbolic embeddings
- Computational efficiency: embeddings quantization, efficient negative sampling

Sequential models

- Advantages

- commonly-used frameworks and optimization algorithms
- easy incremental updates
- sequential patterns learning

- Disadvantages

- no features out-of-the box
- no feature cross
- no user model
- computational limitations

- Useful links

- [The Illustrated GPT-2 \(Visualizing Transformer Language Models\)](#)
- [\[RU\] Интенсив GPT Week Школы анализа данных 2023](#)